



KubeCon



CloudNativeCon

North America 2021

RESILIENCE
REALIZED

Keeping Up with the CVEs!

How to Find a Needle in a Haystack?

Pushkar Joglekar, Sr. Security Engineer, VMware

About Me



- Pushkar Joglekar (he/him)
- Nashik -> Pune -> S.F. Bay Area
- Speaks Marathi, Hindi and English
- Sr. Security Engineer @ VMware Tanzu
- Previously Security Engineer @ Visa (End User)
- Co-wrote “The Kubernetes Book” with Nigel Poulton
- Leads Kubernetes SIG Security Tooling Sub-Project
- CNCF Security TAG Contributor
- Find more about me: <https://pushkarj.github.io/>
- More active on Twitter *than LinkedIn*

What's the deal with the CVEs?



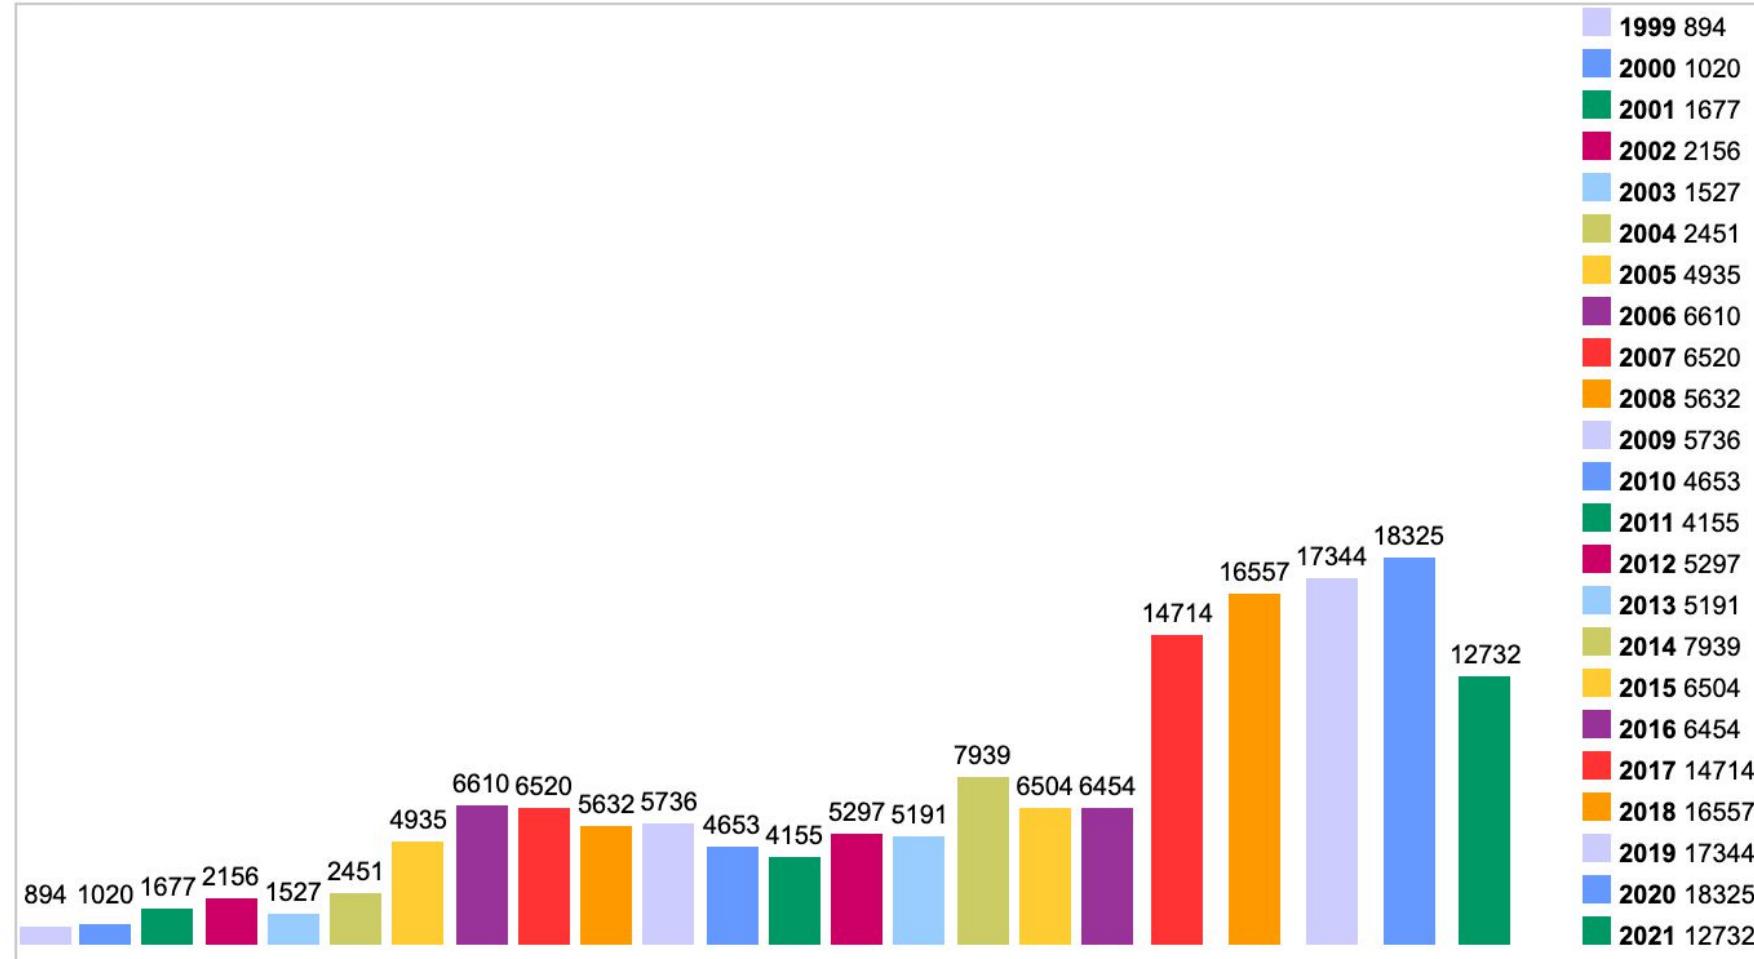
KubeCon



CloudNativeCon

North America 2021

Vulnerabilities By Year

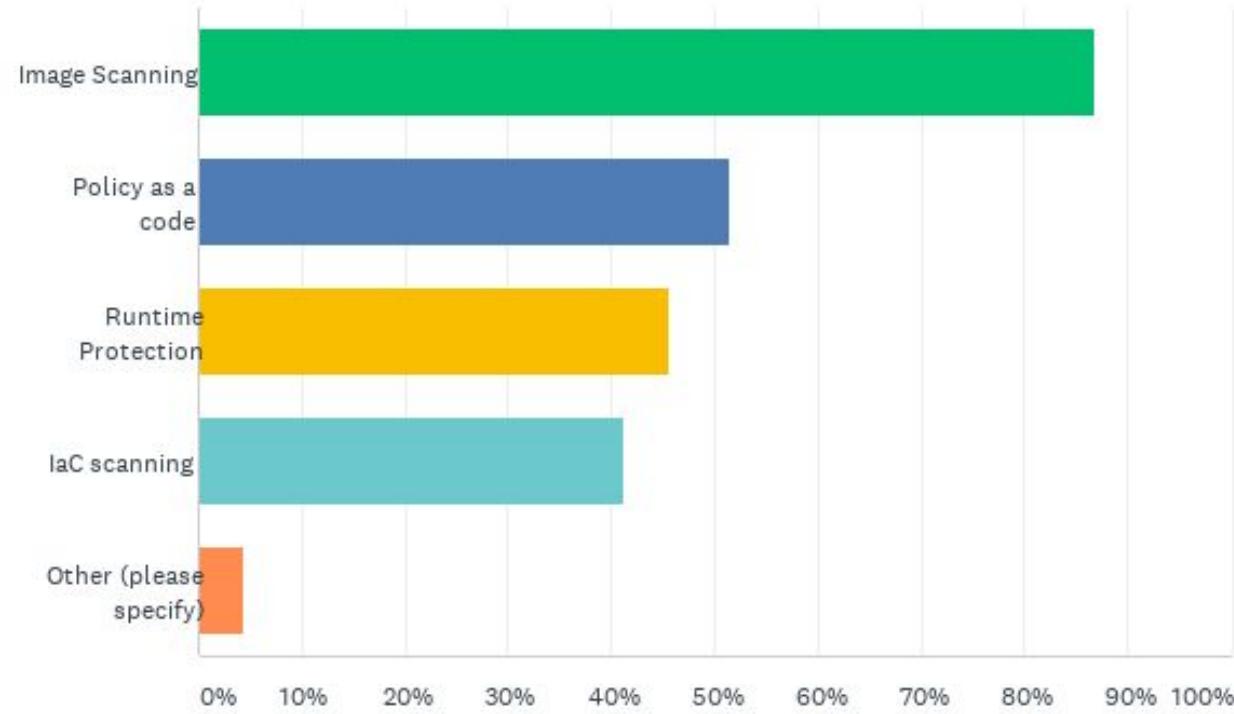


Ref: <https://www.cvedetails.com/browse-by-date.php>



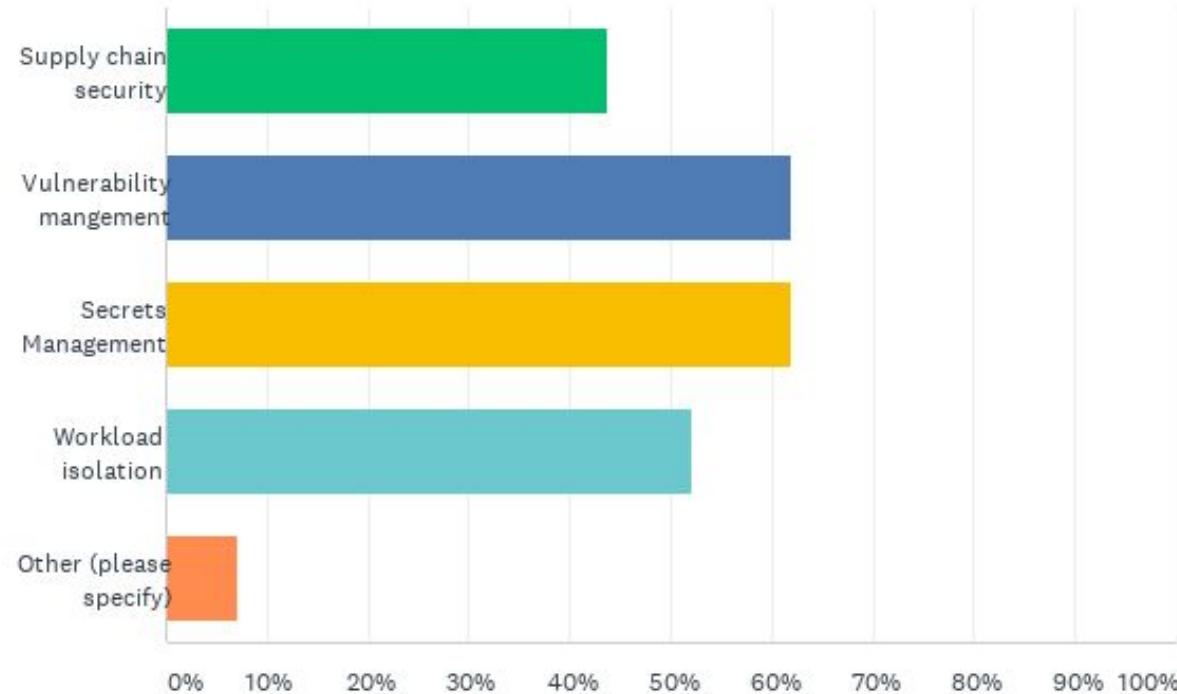
@PuDiJoglekar

Q7 What types of tools are you currently using to protect your cloud native applications?
(Select all that apply)



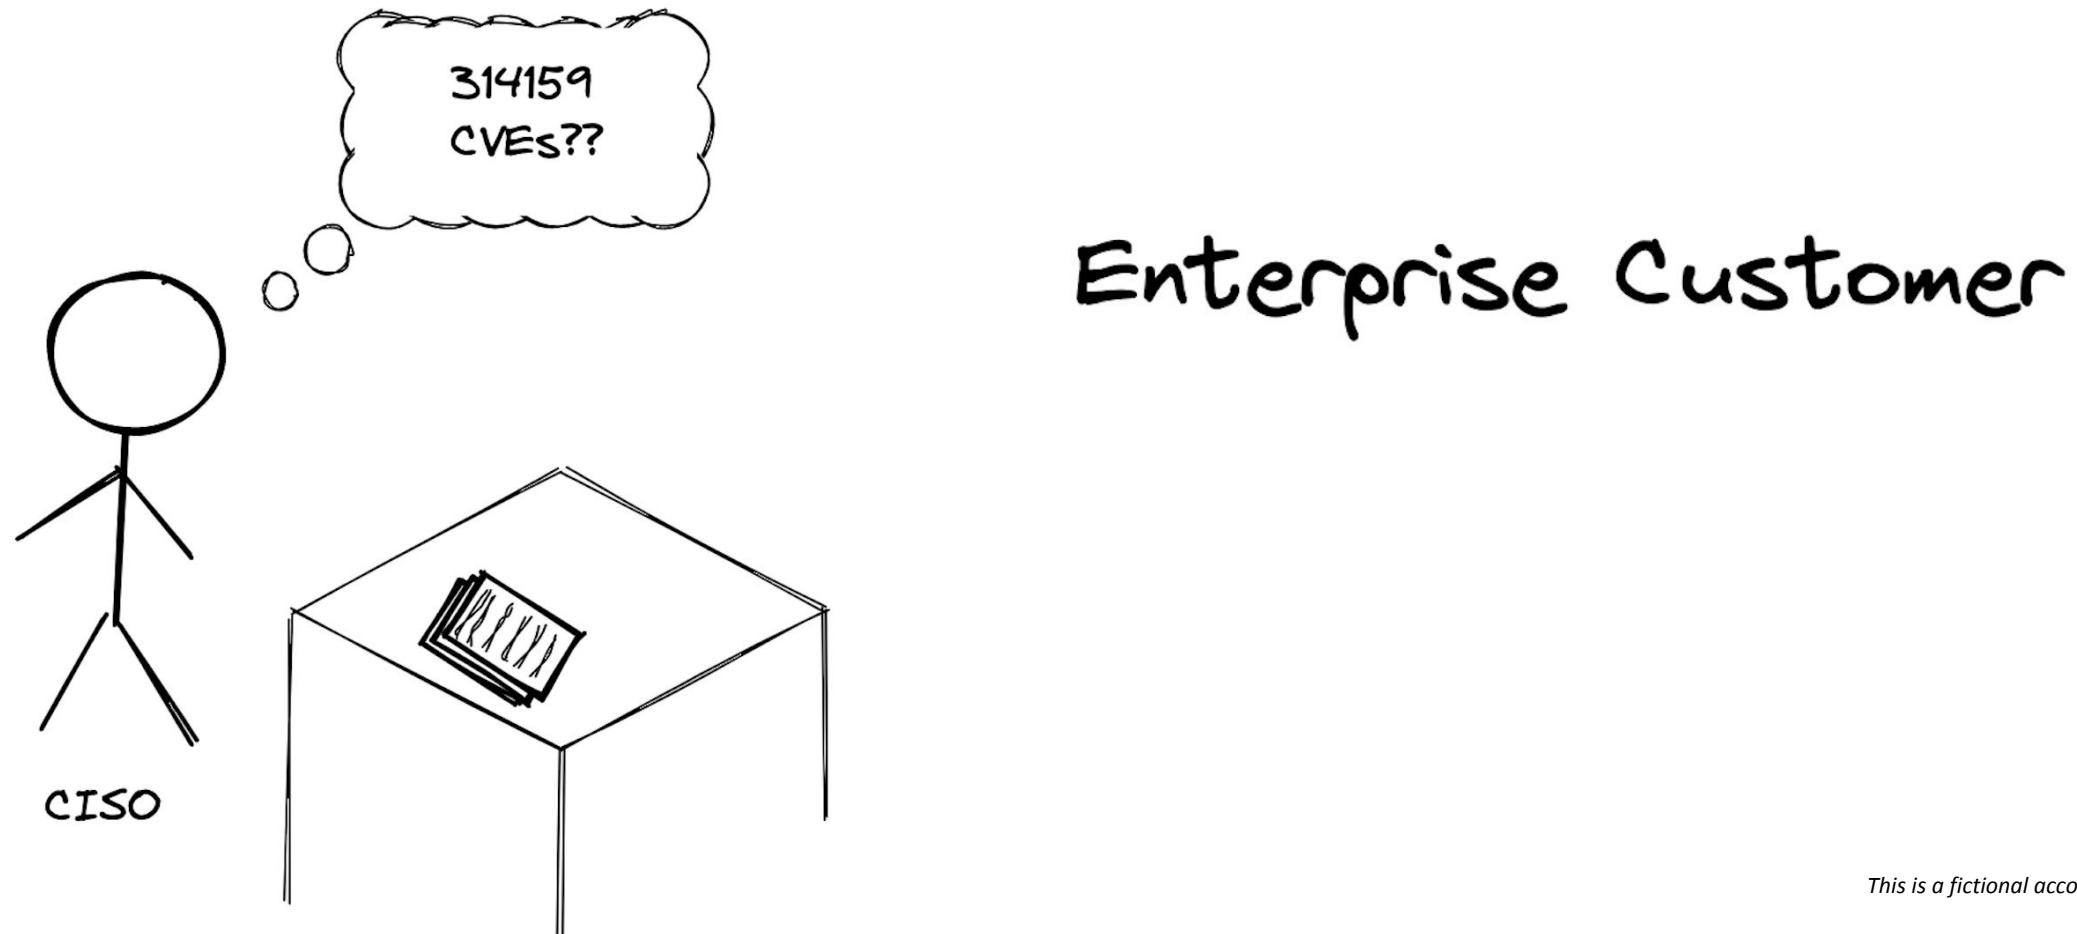
Ref: <https://github.com/cncf/surveys/tree/master/security>

Q4 What is the topmost cloud native security concern for you right now? (Select all that apply)

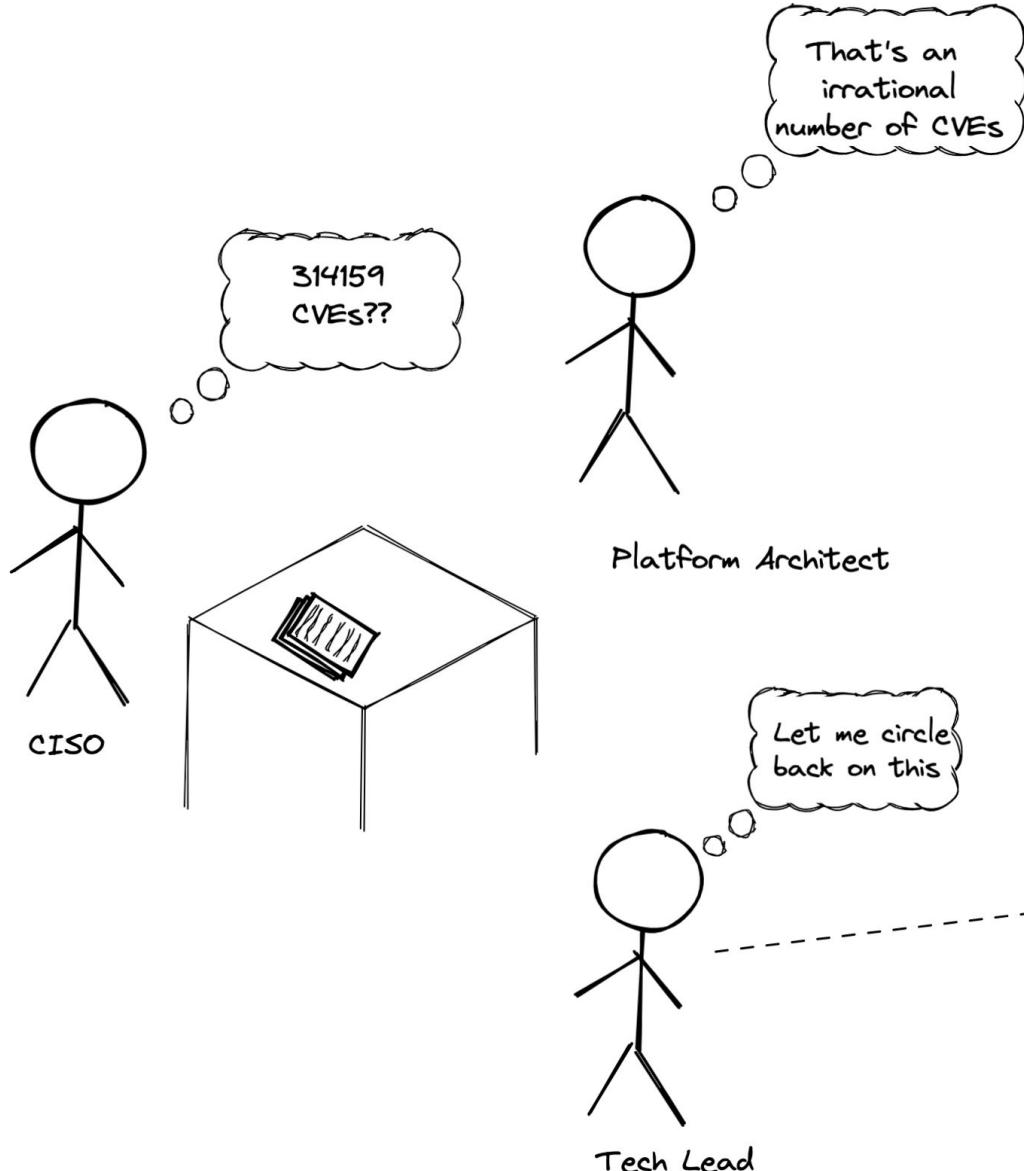


Ref: <https://github.com/cncf/surveys/tree/master/security>

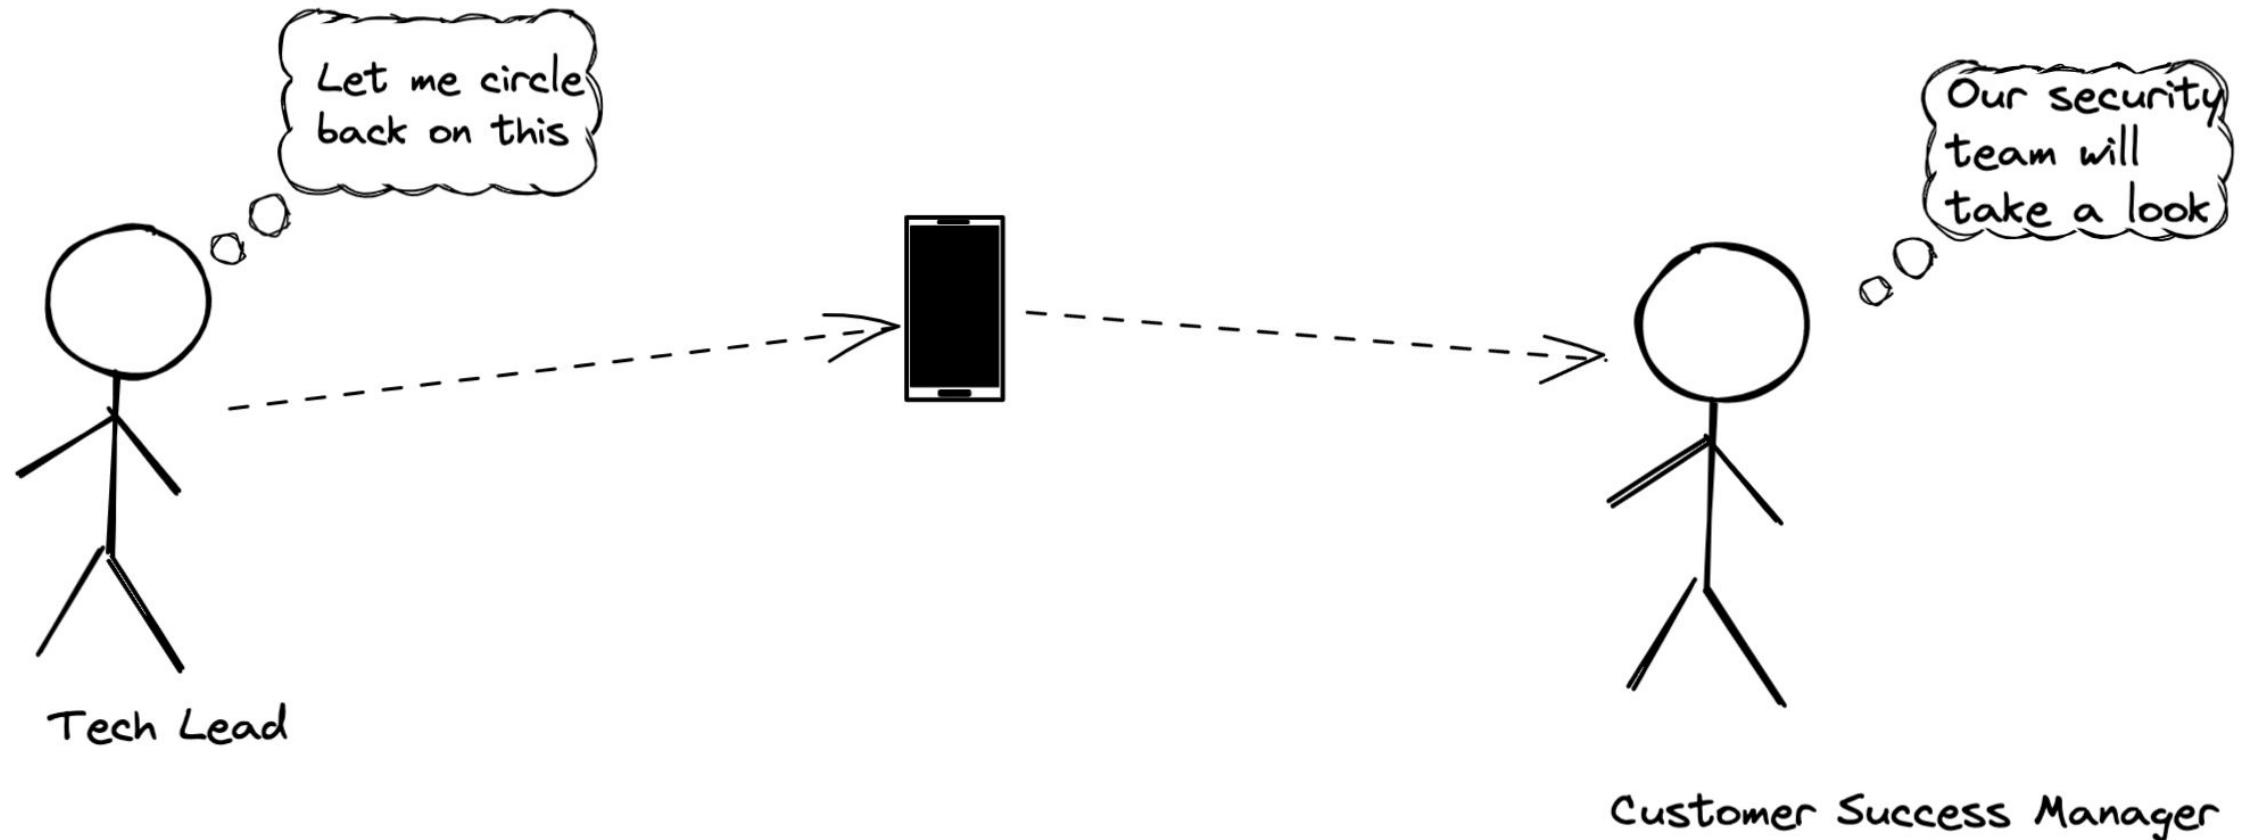
How does this impact real world?



How does this impact real world?



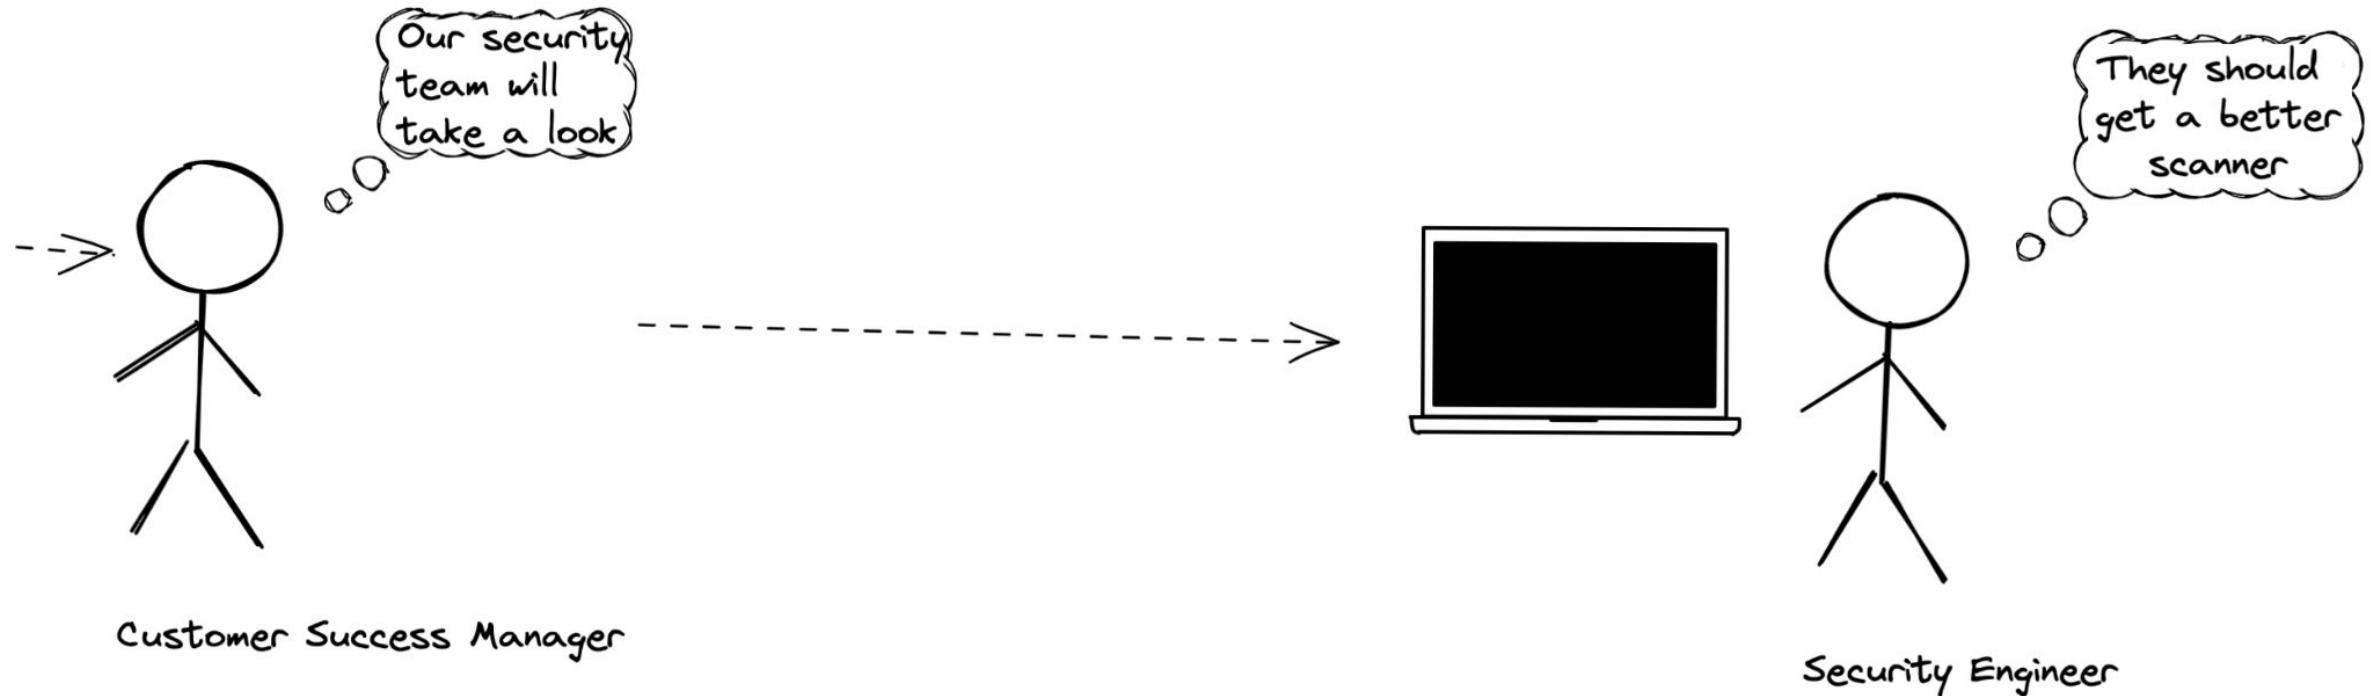
How does this impact real world?



This is a fictional account

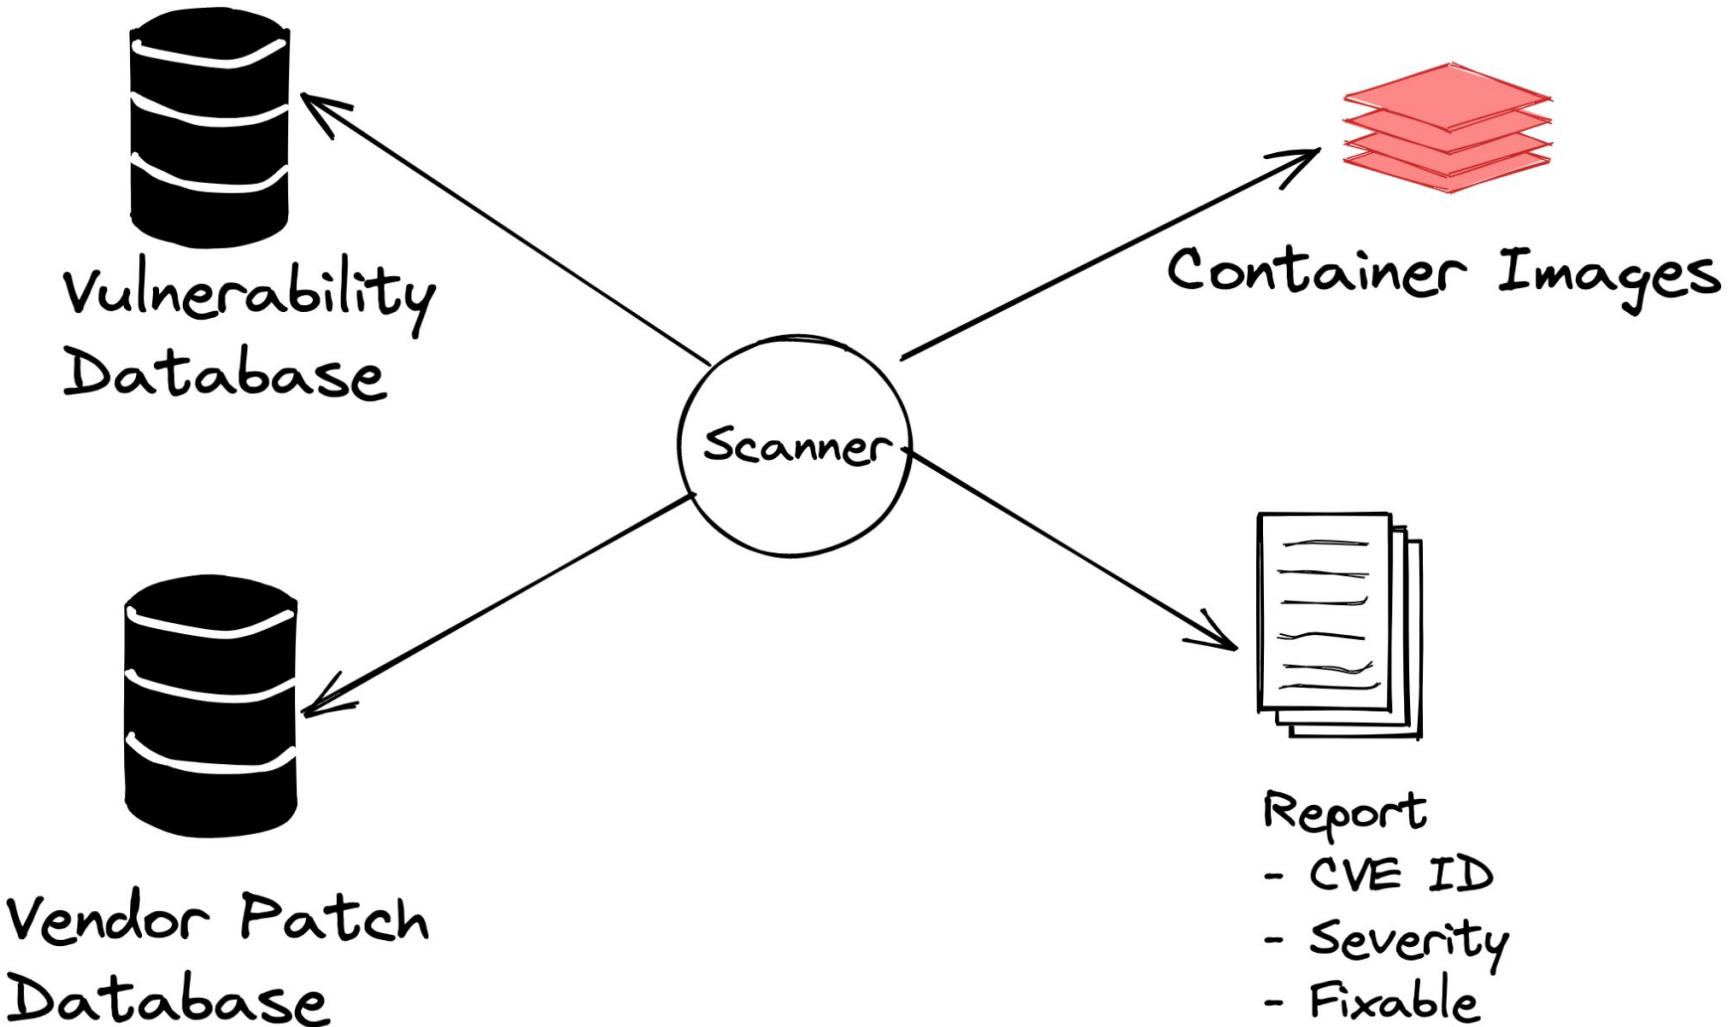
How does this impact real world?

Software Vendor



This is a fictional account

Basic Vulnerability scanner





KubeCon



CloudNativeCon

North America 2021

RESILIENCE
REALIZED

Demo

Analyzing CVE-2019-9923



KubeCon



CloudNativeCon

North America 2021

RESILIENCE
REALIZED

Demo

Analyzing CVE-2018-6829



KubeCon



CloudNativeCon

North America 2021

RESILIENCE
REALIZED

Demo

Analyzing CVE-2021-3711

- Signal to Noise Ratio
 - False positives + negatives
 - Tied to severity than risk
- Impact assessment takes time, effort and expertise
- Risk v/s Reward of updating images
- Too many unneeded packages



Zero Trust

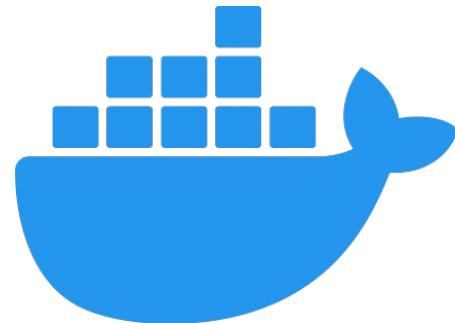
Threat Model

- Knowing the impact of a vulnerability
- Knowing the threat model and risk posture
- Being quick to Detect, Isolate and Fix
- Being able to do all of this in parallel

This is not just an end user problem



kubernetes



docker®

Choosing what to focus on

- Signal to Noise Ratio
 - False positives + negatives
 - Tied to severity than risk
- Impact assessment takes time, effort and expertise
- **Risk v/s Reward of updating images**
- **Too many unneeded packages**

How many images are part of Kubernetes release?

Shoutout to **@Puerco**

```
curl -L https://sbom.k8s.io/v1.22.1/release  
| grep 'PackageName: k8s.gcr.io/'  
| awk '{print $2}'
```



How are these images maintained?

- kube-apiserver
- kube-scheduler
- kube-controller-manager
- kube-proxy
- conformance (skipping)

& etcd

Let's go back to 2020 - Disclaimer

- Many community contributors made this happen!
- Focus is on container images only
- Future is even more exciting!

Let's go back to 2020 - KEPs

Shoutout to **@yuvemna**

Motivation

Rebasing the k8s images to distroless/static can make the images thinner, safer and less vulnerable.

Meanwhile, it will drastically reduce churn on the total number of k8s images versions. Due to the fact that many images are based on debian base and a vulnerability in debian base (a couple times a month) will result in rebuilding every image, changing the image from debian base to distroless/static can reduce the total number of k8s image versions.

What's more, it reduces the burden of managing and maintaining multiple k8s images from the security (e.g. CVE), compatibility and build process concerns.

Ref: <https://github.com/kubernetes/enhancements/tree/master/keps/sig-release/1729-rebase-images-to-distroless>

But first what is distroless?



WH-WHAT IS A WEEK-END?



Distroless: Base & Static

Statically compiled applications (Go) that do not require libc can use the `gcr.io/distroless/static` image, which contains:

- ca-certificates
- A /etc/passwd entry for a root user
- A /tmp directory
- tzdata

Most other applications (and Go apps that require libc/cgo) should start with `gcr.io/distroless/base`, which contains all of the packages in `gcr.io/distroless/static`, and

- glibc
- libssl
- openssl

Ref: <https://github.com/GoogleContainerTools/distroless/tree/main/base>

Security v/s Usability



Before

```
"/bin/sh",
  "-c",
  "exec /usr/local/bin/kube-apiserver </snip> 1>>/var/log/kube-apiserver.log
2>&1"
```

After

```
"/go-runner", "--log-file=/var/log/kube-apiserver.log", "--also-stdout=false",
"--redirect-stderr=true",
"/usr/local/bin/kube-apiserver",
</snip>
```

Ref: <https://github.com/kubernetes/kubernetes/pull/90804>

Switch to static bash and distroless image for etcd #91171

Merged

k8s-ci-robot merged 2 commits into [kubernetes:master](#) from [dims:switch-etcd-to-bash-static](#) on May 19, 2020

Conversation 39

Commits 2

Checks 0

Files changed 14



dims commented on May 16, 2020 · edited

Member  ...

What type of PR is this?

/kind feature

What this PR does / why we need it:

This is similar to [#90674](#) where we switched apiserver/scheduler to distroless. In this case etcd image has scripts, the idea is to use the `bash-static` utility (instead of the FULL debian-base image) inside the distroless image and cleanup a bunch of things to get this to work:



@PuDiJoglekar



TL;DR kube-proxy needed IPtables

So it had to continue to use Debian upstream image.

As a result it is now the image that is most often updated :-)

What did we learn from this?



Release train rarely looks like this



It can often look like this



But an imperfect solution still has benefits

```
pjoglekar@pjoglekar-a01 demo % trivy image
gcr.io/google-containers/kube-apiserver:v1.16.9

gcr.io/google-containers/kube-apiserver:v1.16.9 (debian 9.12)
=====
Total: 167 (UNKNOWN: 1, LOW: 93, MEDIUM: 26, HIGH: 28, CRITICAL: 19)
```

```
pjoglekar@pjoglekar-a01 demo % trivy image
k8s.gcr.io/kube-apiserver:v1.22.1

k8s.gcr.io/kube-apiserver:v1.22.1 (debian 10.10)
=====
Total: 0 (UNKNOWN: 0, LOW: 0, MEDIUM: 0, HIGH: 0, CRITICAL: 0)
```

Ref: Snippet from <https://github.com/aquasecurity/trivy> scan result

None of the vulnerable
components were needed



Imagine the reduced churn in fixing these vulnerabilities

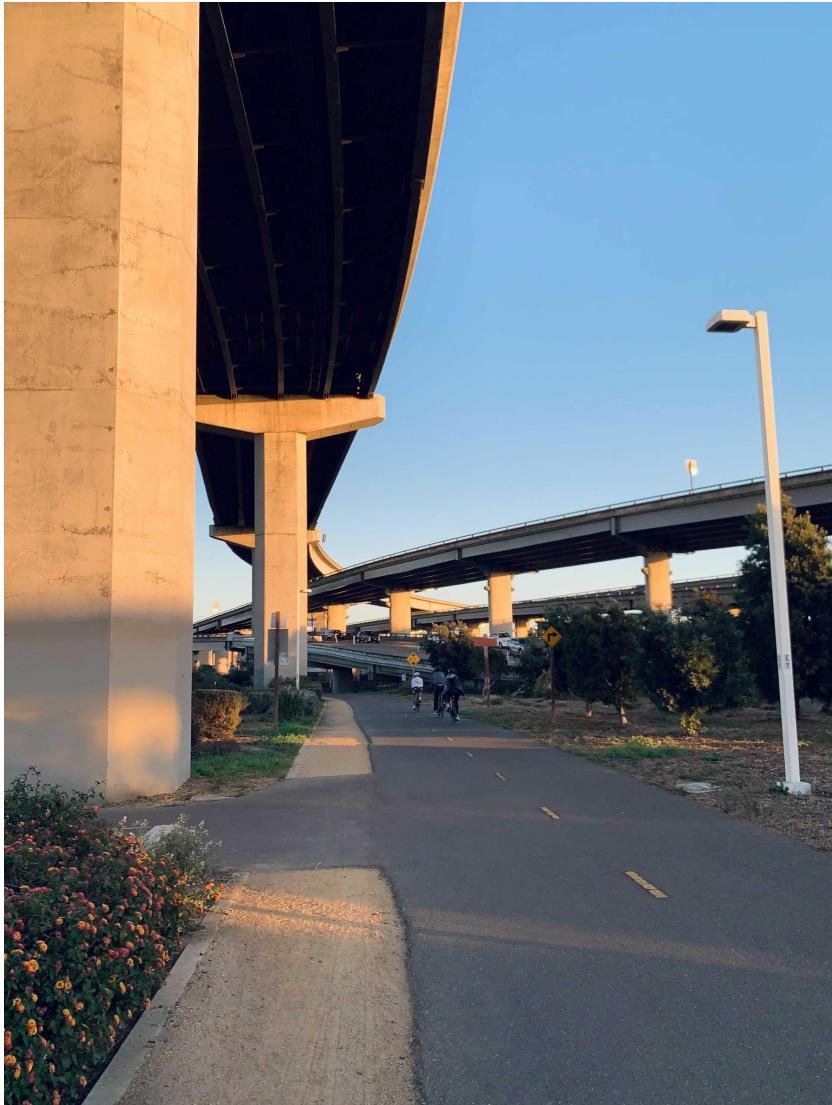


What to do when distroless is not an option?



Focus on fixable CVEs

What to do when distroless is not an option?



Focus on whether the vulnerability
is in code execution path or not

What to do when distroless is not an option?



Focus on improving automation to rebuild, ship and test updated base images

What to do when distroless is not an option?



Create a list of images that
need special care / frequent
updates e.g. kube-proxy

- Defining a policy and process to triage **known** vulnerabilities
- Automated Periodic scans for **known** vulnerabilities:
 - ◆ In Images and
 - ◆ Build time dependencies
- Feedback or Suggestions welcome!!
 - ◆ [#sig-security-tooling](#) on Kubernetes Slack

Manage Vulnerabilities, by *Being* Vulnerable

- Identify base images and standardize them
- Learn limitations of vulnerability scanners
- Understand the efforts to remediate

... Most importantly model threats and assess risk

Questions?



@PuDiJoglekar

Thank you!



@PuDiJoglekar